

Computer Architecture (DAT 105 / DIT 051)

Credit points: 7.5 EHC

Course Examiner and Lecturer: Per Stenstrom

Teaching Assistant:

- Fazeleh Sadat Hoseini
- Piyumal Ranawaka
- Qi Shao

Department of Computer Science & Engineering
Chalmers University of Technology

What is Computer Architecture All About?




A Fascinating Success Story

Intel i9 (2017)
~ 30 billion ops/s,
~ 7 billion transistors

ENIAC (1946):
~ 1000 ops/s 18000 vacuum tubes
~ 10 m



**Millions of times faster,
smaller, and more cost effective**



Moore's Law

Number of transistors
(logarithmic)

1000 000 000

Twice as many
transistors about
every 2 years!
Now tapering off to
every 3-4 years

1000

1970

2015


- How can we tame all these transistors to compute?

CHALMERS

Chalmers University of Technology



Computer Architecture



The engineering discipline of computer design

The hardware/software interface

Instruction Set Architecture (ISA)

Computer organization

Hardware design



Quiz 1.1

Compare ENIAC with Intel Core I9. Which of the following statements are true

- ☒ A. Intel I9 is at least 20 million times faster than ENIAC
- ☐ B. Intel I9 is at most 2 million times faster than ENIAC
- ☒ C. Core I9 contains at least 300 thousand more switching elements than ENIAC
- ☒ D. Core I9 is at least 70000 times smaller than ENIAC
- ☐ E. Core I9 is at least 1 million times smaller than ENIAC



How is the Course Organized?

Learning Objectives

After the course you should:

- **master** fundamental **concepts and terminology**
- **understand design principles of processors**
- **understand design principles of memory hierarchies**
- **understand design principles of multicore microprocessors**
- **be able to use modeling methods** to assess the impact design alternatives have on **performance / energy**


Course Organization

- 7 lectures & 2 optional crash-course lectures
- 2 real-stuff presentations
- 6 problem-solving sessions
- 3 laboratory sessions/assignments (reports)

Lectures: Flipped-classroom teaching. Interaction. Active learning!

Preparation: Watch available youtube lectures before class (starting this week)!!!!

Textbook



- Book is available for sale at Cremona
- Only selected chapters of the book
- Book is also used in other courses in the Computer Architecture track.

Course Contents – Lectures

1. **Introduction:** Components of a computer, technology trends, performance principles and metrics

Opt 1 (Basics of modern processor design: **Pipelining concepts**)

*Do you know what a **RAW hazard** is? **Delayed branch**? **Stalling**?*

Opt 2 (Basics of memory hierarchies: **Caches & Interconnects**)

*Do you know what a **memory hierarchy** is? **Set associativity**? **Miss penalty**?*

- 2-4. **Techniques to uncover ILP:** Overcoming limitations of instruction ordering, name dependencies, and control flow: Static instruction scheduling, out-of-order execution, branch prediction, speculative execution
5. **Memory hierarchy concepts.** Cache and virtual memory. The #C model, inclusion, lockup-free caches, prefetching etc.
6. **Techniques to uncover TLP:** Thread-level parallelism, multicore-multithreaded architectures, cache coherence
7. **Very Large Instruction Word (VLIW) computers.** Static techniques for uncovering ILP, trace scheduling, predicated execution.

Course Schedule

TimeEdit

Course information:

- At Canvas
- Course Schedule (under Files/Syllabus tab)
- Course material (under Files tab)
 - Video lectures (links to Youtube and slides)
 - Flipped classroom material (slides)
 - Exercises
 - Laboratory PMs

Keep an eye on Announcements at Canvas regularly!!!

- Real-time information about course logistics

Examination

- **Laboratory project assignment**
 - Brief written report required for each assignment
- **Quizzes at each lecture**
 - Based on youtube lectures
 - 3 correct => **4 bonus points for the exam** (for higher grades only)
 - Bring your laptop; we will use Canvas for quizzes
- **Real-stuff studies**
 - Learn how principles are used in commercial machines. **Another 4 bonus points** (for higher grades only)
- **Written exam**
 - Typically on concepts and problem solving. Examples of old exams will be distributed. **Date: October 24, 2022.**

Course Evaluation

- **We want to improve the course continuously**
 - Does the course start at the right level?
 - How good is the teaching material (textbook, problems, lab manuals)?
 - How well do the instructors perform?
 - What is the pace in the course?

Course representatives:

- FK Benjamin Björklund
- MPHPC Long Cheng
- MPHPC Tahmina Aziz Tahrin

Please give us feedback; we use it to improve the course – in real-time, and for coming generations of students



Michel Dubois, Murali Annavaram, Per Stenström © 2017

Lecture 1

Fundamentals of Computer Architecture



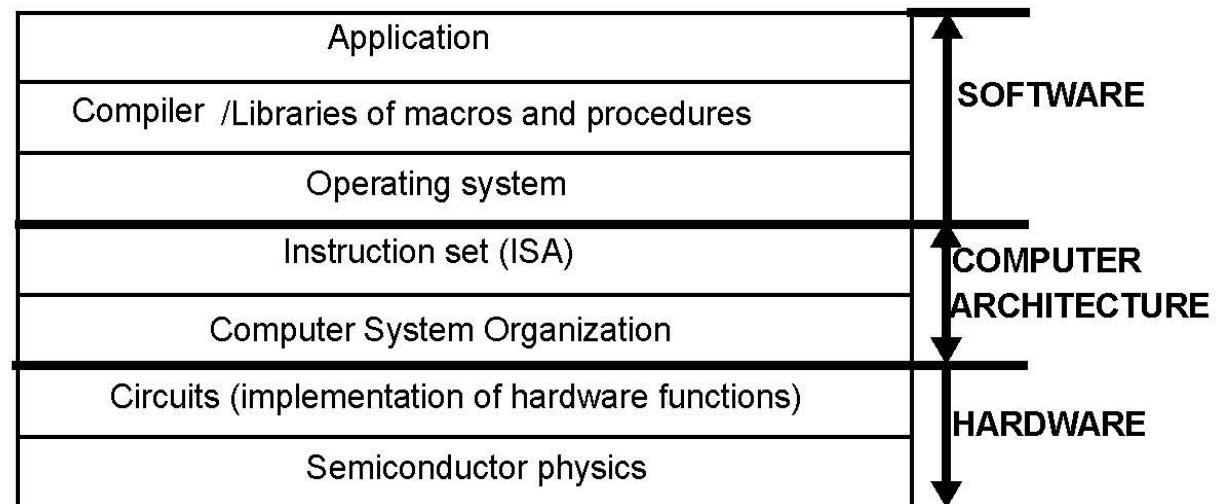
Components of a Computer, Trends and Parallelism (Ch 1.1 – 1.3, 1.5)



What is Computer Architecture?

- **Old definition:** Instruction Set Architecture (ISA)
- **Today's definition:** Hardware organization of computers (including ISA)

The Transformation Hierarchy:




The role of a computer architect: To make design trade-offs across the HW/SW interface to meet functional, performance and cost requirements

A PC Computer Organization

The organization of a modern PC:

Processor: Today multicore

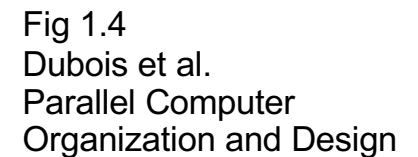


Memory

Interconnect

I/O

Fig 1.3
Dubois et al.
Parallel Computer
Organization and Design



Processor (P), memory system (M and C), I/O, and networks

Processor Architecture 1(2)

- Clock rates of microprocessors have increased exponentially





Fig 1.5
Dubois et al.
Parallel Computer
Organization and Design

Key factors:

- Technology improvements only (19% annual improv.)
Deeper pipelines + architecture innovations
- Altogether 49% up until 2002

Processor Architecture 2(2)

Computer architects take advantage of the growing number of transistors:



- New process every two year. Feature size reduced by 30% in every new process
- Transistor count doubles every 2 years (Moore's law)

- A sandbox to play in so to speak
- How do we use 100B transistors?

This trend is tapering off making architecture improvements the more important

Memory Systems

- Main memory speed grows slower than processor speed (“Memory wall”)
- A multi-level memory hierarchy tries to realize a memory that is big, fast and cheap
- It works surprisingly well due to the **Principle of locality**

Memory	Size	Marginal Cost	Cost per MB	Access time
L2 Cache(on chip)	1MB	\$20/MB	\$20	5nsec
Main Memory	1GB	\$50/GB	5c	200nsec
Disk	500GB	\$100/500GB	0.02c	5msec


Table shows the cost and speed of memories in a PC (2008)

The Memory Wall

Microprocessor speed has increased by 50% every year

- DRAM performance has improved by only 7% per year
- DRAM density keeps increasing by 4X every 3 years

Trends have changed dramatically in the last years



DRAM: 1.07 CGR

Memory wall =
 $\text{memory_cycle} / \text{processor_cycle}$

- In 1990: 4 (25MHz, 150ns)
- In 2002: ~ 200
- > 2002: Is tapering off/reversing

Multicores have turned memory wall into a bandwidth wall

CHALMERS

Chalmers University of Technology

Parallelism in Architectures 1(2)

Scalar processors do operations on single (not vector) operands

- A typical scalar instruction

Example) ADD O1,O2,O3 / $O2+O3 \Rightarrow O1$

Multiple scalar instructions can be executed at a time:

- Pipelining, superscalars, superpipelining

These approaches exploit ***Instruction-Level Parallelism (ILP)*** exposed in a single thread/process execution

Chip multiprocessors (or multicores) exploit parallelism across threads running in parallel on different processors/cores: (***Thread-Level Parallelism TLP***)



Parallelism in Architectures 2(2)

Vector and array processors



- Example vector instruction

VADD VO1,VO2,VO3 /VO2+VO3==>VO1

- VO_k is a vector of scalar components
- Equivalent to computing

VO2[i]+VO3[i]==>VO1[i], i=0,1,...,N

Vectors are executed by pipelines or parallel arrays



Performance Metrics and Evaluation (Ch 1.4)



Performance Metrics (Measures)

Metric #1: Time to complete a task (T_{exe}): Execution time, response time, latency

- “X is N times faster than Y” means $T_{exe}(Y)/T_{exe}(X) = N$
- The major metric used in this course

Metric #2: Number of tasks per time unit (e.g., hr, s, ns)

- The throughput for X is N times higher than Y if
 $Throughput(X)/Throughput(Y) = N$

NOTE: Does not always lead to same design options

Example of unreliable metrics: MIPS and MFLOPS

**Execution time is the ultimate measure of performance
=>Benchmarking**

Which Programs to Choose?

Early unsuccessful attempts

Real programs:

- **Problems:** Porting, complexity and not easy to interpret results

Kernels

- Computationally intense piece of real program
- **Problem:** Does not capture impact of entire program

Toy benchmarks:

- **Problem:** Unclear link to reality

Synthetic benchmarks

- **Problem:** Inherently unreal

Benchmark Suites

Commonly used benchmark suites

- SPEC: Standard Performance Evaluation Corporation
- Scientific/Engineering/General Purpose
- Integer and floating point intense program suites
- New editions in 1995, 1998, 2000, 2006 and 2017
- TPC benchmarks (transaction processing)
- Embedded benchmarks (EEMBC)
- Media benchmarks

Reporting Performance

Given a set of N programs

(Weighted) arithmetic mean: $\sum_i T_i / N$ **or** $\sum_i T_i \times W_i$

Speedup over a reference machine R: $S_i = \frac{T_{R,i}}{T_i}$

Geometric mean of speedup: $\bar{S} = \sqrt[N]{\prod_{i=1}^N S_i}$

Problem with Arithmetic Mean (Lower is better)

	Program 1 [s]	Program 2 [s]	Program 3 [s]
Machine 1	1	2	10
Machine 2	2	3	4
Reference	10	24	60

	Arithmetic mean [s]
Machine 1	4.3
Machine 2	3
Reference	31.3

Outliers!

Geometric Mean (Higher is better)

	Program 1 [s]	Program 2 [s]	Program 3 [s]
Machine 1	1	2	10
Machine 2	2	3	4
Reference	10	24	60

	S for P1	S for P2	S for P3	Geom. mean
Machine 1	$10/1 = 10$	$24/2 = 12$	$60/10 = 6$	8.96
Machine 2	$10/2 = 5$	$24/3 = 8$	$60/4 = 15$	8.43

Reference

Used in performance reports (e.g., SPEC

CHALMERS

Chalmers University of Technology

Quiz 1.2

The execution times for programs P1, P2 and P3: 1, 4 and 8 s on machine M1 and 2, 8, 16 s on a reference machine R.

Which of the following statements are correct?

- A) The average execution time of M1 is 4.3 s
- B) The average execution time of R is 8.6 s
- C) M1 is $(8.6/4.3 =) 2 \times$ faster than R
- D) R is $(8.6/4.3 =) 2 \times$ faster than M1
- E) Using R as a reference, the geometric mean of speedup is 2


Fundamental Performance Equations for CPUs

$$T_{exe} = IC \times CPI \times T_c$$

- **IC:** Depends on program, compiler and ISA.
- **CPI:** Depends on *instruction mix*, ISA, and implementation
- **T_c:** Depends on implementation and technology

When processor executes more than one instruction per clock use: **$T_{exe} = (IC \times T_c)/IPC$**

Amdahl's Law



Enhancement ***E*** accelerates a fraction ***F*** of the task by a factor ***S***

$$T_{\text{exe}}(\text{with } E) = T_{\text{exe}}(\text{without } E) \times \left[(1 - F) + \frac{F}{S} \right]$$

$$\text{Speedup}(E) = \frac{T_{\text{exe}}(\text{without } E)}{T_{\text{exe}}(\text{with } E)} = \frac{1}{(1 - F) + \frac{F}{S}}$$

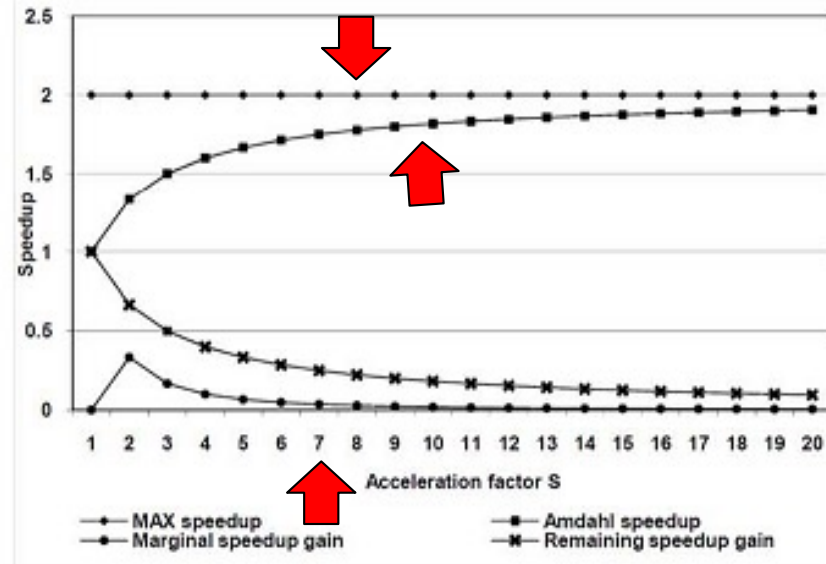
Lessons Learned from Amdahl's Law

LESSON #1: Improvement is limited by the fraction of the execution time that cannot be enhanced:

$$\text{SPEEDUP}(E) < \frac{1}{1-F}$$

LESSON #2: Optimize the common case; for example, deal with rare cases in software.

LESSON #3: Law of diminishing returns: Gains diminish as we add more resources:



F=0.5



Power 1(2)

Total Power = Dynamic + Static (leakage)

$$P_{dynamic} = \alpha C V^2 f$$

$$P_{static} = V I_{sub} \propto V e^{-k \frac{V_T}{T}}$$

Dynamic power favors parallel processing over higher clock frequency. Dynamic power roughly proportional to frequency³

Static power: Circuits leak independently from frequency

Power 2(2)

Power/energy are critical problems

- Power (immediate energy dissipation) must be dissipated
- Otherwise, temperature goes up which affects performance, correctness, and may eventually destroy the circuit
- Effect on the supply of power to the chip
- Energy (depends on power and speed)
- Costly; a global problem
- Problematic for battery-operated devices



Quiz 1.3

The execution time of a program on a single processor is 100s. 50% of that execution has a speedup linear to the number of processors

Which of the following statements are correct?

- ☒ A) The execution time of the program on two processors is 50 s
- ☒ B) The execution time of the program on two processors is 75 s
- ☒ C) The execution time of the program can never go below 50 s regardless of the number of processors
- ☒ D) The execution time of the program with 100 processors is 1 s

What you should know by now

- What is computer architecture all about and what is the role of a computer architect
- Components of a computer organization and how trends affect them (Moore's law, Memory wall...)
- Parallelism in computer architecture: ILP and TLP
- Quantitative metrics and methods for reporting performance:
 - Benchmarking
 - Execution time and speedup (arithmetic, harmonic, geometric means)
 - Amdahl's Law